

Symbolic Machine Learning

The main way-to-go for implementing a machine learning pipeline in Julia is via the [MLJ.jl](#) package.

We are going to work with the `iris` dataset, trying to discover the relation between the specific attribute values of an iris flower and the family to which the same flower belongs to.

More generally, we want to find the relation between the values of the *attributes* of each instance (`X`) and the corresponding *labels* (`y`).

In order to do so, we are going to train a (classification) decision tree, leveraging the `DecisionTree` package, which can be easily integrated within an `MLJ` pipeline.

Later in this notebook, we will repeat this process leveraging the `Sole.jl` library, which will allow us to explicitly model the problem through the lens of logic.

```
In [1]: using Pkg  
Pkg.activate("..")  
Pkg.instantiate()  
Pkg.update()  
  
Activating project at `~/logic-and-machine-learning`  
Updating registry at `~/.julia/registries/General.toml`  
Updating git-repo `https://github.com/aclai-lab/ManyExpertDecisionTree  
s.jl`  
Updating git-repo `https://github.com/aclai-lab/SoleReasoners.jl#embed  
ding`  
No Changes to `~/logic-and-machine-learning/Project.toml`  
No Changes to `~/logic-and-machine-learning/Manifest.toml`  
  
In [2]: # for reproducibility purposes  
using Random  
Random.seed!(1605)  
  
Out[2]: TaskLocalRNG()
```

Learning with MLJ.jl

Data Loading and Description

```
In [3]: using MLJ  
using RDatasets # used to load the iris dataset  
  
data = RDatasets.dataset("datasets", "iris");  
  
In [4]: schema(data)
```

Out[4]:

names	scitypes	types
SepalLength	Continuous	Float64
SepalWidth	Continuous	Float64
PetalLength	Continuous	Float64
PetalWidth	Continuous	Float64
Species	Multiclass{3}	CategoricalValue{String, UInt8}

In [5]: `data`

Out[5]: 150×5 DataFrame

125 rows omitted

Row	SepalLength	SepalWidth	PetalLength	PetalWidth	Species
	Float64	Float64	Float64	Float64	Cat...
1	5.1	3.5	1.4	0.2	setosa
2	4.9	3.0	1.4	0.2	setosa
3	4.7	3.2	1.3	0.2	setosa
4	4.6	3.1	1.5	0.2	setosa
5	5.0	3.6	1.4	0.2	setosa
6	5.4	3.9	1.7	0.4	setosa
7	4.6	3.4	1.4	0.3	setosa
8	5.0	3.4	1.5	0.2	setosa
9	4.4	2.9	1.4	0.2	setosa
10	4.9	3.1	1.5	0.1	setosa
11	5.4	3.7	1.5	0.2	setosa
12	4.8	3.4	1.6	0.2	setosa
13	4.8	3.0	1.4	0.1	setosa
:	:	:	:	:	:
139	6.0	3.0	4.8	1.8	virginica
140	6.9	3.1	5.4	2.1	virginica
141	6.7	3.1	5.6	2.4	virginica
142	6.9	3.1	5.1	2.3	virginica
143	5.8	2.7	5.1	1.9	virginica
144	6.8	3.2	5.9	2.3	virginica
145	6.7	3.3	5.7	2.5	virginica
146	6.7	3.0	5.2	2.3	virginica
147	6.3	2.5	5.0	1.9	virginica
148	6.5	3.0	5.2	2.0	virginica
149	6.2	3.4	5.4	2.3	virginica
150	5.9	3.0	5.1	1.8	virginica

In [6]: `y, X = unpack(data, ==(:Species))`

Row	SepalLength	SepalWidth	PetalLength	PetalWidth
	Float64	Float64	Float64	Float64
1	5.1	3.5	1.4	0.2
2	4.9	3.0	1.4	0.2
3	4.7	3.2	1.3	0.2
4	4.6	3.1	1.5	0.2
5	5.0	3.6	1.4	0.2
6	5.4	3.9	1.7	0.4
7	4.6	3.4	1.4	0.3
8	5.0	3.4	1.5	0.2
9	4.4	2.9	1.4	0.2
10	4.9	3.1	1.5	0.1
11	5.4	3.7	1.5	0.2
:	:	:	:	:
141	6.7	3.1	5.6	2.4
142	6.9	3.1	5.1	2.3
143	5.8	2.7	5.1	1.9
144	6.8	3.2	5.9	2.3
145	6.7	3.3	5.7	2.5
146	6.7	3.0	5.2	2.3
147	6.3	2.5	5.0	1.9
148	6.5	3.0	5.2	2.0
149	6.2	3.4	5.4	2.3
150	5.9	3.0	5.1	1.8

```
In [7]: # categorical vectors are lighter than raw vectors; can you guess why?  
typeof(y)
```

```
Out[7]: CategoricalVector{String, UInt8, String, CategoricalValue{String, UInt8}, Union{}}
          (alias for CategoricalArrays.CategoricalArray{String, 1, UInt8, String, CategoricalArrays.CategoricalValue{String, UInt8}, Union{}})
```

In [8]: `typeof(X)`

Out[8]: DataFrame

```
In [9]: # to ensure that classes are balanced
```

```
for class in unique(y)
    println("${class} - ${count(yi -> yi == class, y))")
end
```

setosa - 50
versicolor - 50
virginica - 50

Data Preprocessing

In the limited scenario of this exercise, there is not much space for complex preprocessing of our data. For example, we are not dealing with unbalanced classes, missing data, or complex encodings.

The usual workflow, at this point, is to partition the data into a training and a test bucket, keeping a balanced class diversity.

With this distinction, we can train a model on the initial training data and leverage the test one for simulating a real-world scenario, obtaining reliable performance.

MLJ makes our work *much* easier, even providing us with a more sophisticated training strategy, as we will see later.

Model Training

We will integrate an external model, coming from the `DecisionTree` package, into the MLJ workflow.

In the next lessons, we will be doing something similar with another model called `ModalDecisionTree`.

```
In [10]: try
    DecisionTreeClassifier = @load DecisionTreeClassifier pkg=DecisionTree
catch
    println("The DecisionTreeClassifier symbol has already been imported.
end
```

```
[ Info: For silent loading, specify `verbosity=0`.
import MLJDecisionTreeInterface ✓
```

```
Out[10]: MLJDecisionTreeInterface.DecisionTreeClassifier
```

```
In [11]: model = MLJDecisionTreeInterface.DecisionTreeClassifier(
    max_depth=5,
    min_samples_leaf=1,
    min_samples_split=2
)
```

```
Out[11]: DecisionTreeClassifier(
    max_depth = 5,
    min_samples_leaf = 1,
    min_samples_split = 2,
    min_purity_increase = 0.0,
    n_subfeatures = 0,
    post_prune = false,
    merge_purity_threshold = 1.0,
    display_depth = 5,
    feature_importance = :impurity,
    rng = TaskLocalRNG())
```

A machine is a binding between a model and the data it works with.

It also keeps track of other information we might want to inspect, such as the specific parameters learned by a model.

In the cell below, we bind the decision tree model to all the instances we have available. This is not a good idea, but we will return on the topic in a moment.

```
In [12]: mach = machine(model, X, y)
```

```
Out[12]: untrained Machine; caches model-specific representations of data
          model: DecisionTreeClassifier(max_depth = 5, ...)
          args:
            1: Source @299 ↴ Table{AbstractVector{Continuous}}
            2: Source @769 ↴ AbstractVector{Multiclass{3}}
```

```
In [13]: fit!(mach)
```

```
[ Info: Training machine(DecisionTreeClassifier(max_depth = 5, ...), ...).
```

```
Out[13]: trained Machine; caches model-specific representations of data
          model: DecisionTreeClassifier(max_depth = 5, ...)
          args:
            1: Source @299 ↴ Table{AbstractVector{Continuous}}
            2: Source @769 ↴ AbstractVector{Multiclass{3}}
```

```
In [14]: y_predict_probabilities = MLJ.predict(mach, X)
          y_predict = mode.(y_predict_probabilities)
```

```
Out[14]: 150-element CategoricalArrays.CategoricalArray{String,1,UInt8}:
  "setosa"
  "virginica"
  "virginica"
```

```
In [15]: fitted_params(mach).tree
```

```
Out[15]: PetalLength < 2.45
    └─ setosa (50/50)
        └─ PetalWidth < 1.75
            └─ PetalLength < 4.95
                └─ PetalWidth < 1.65
                    └─ versicolor (47/47)
                        └─ virginica (1/1)
                └─ PetalWidth < 1.55
                    └─ virginica (3/3)
                └─ SepalLength < 6.95
                    └─ versicolor (2/2)
                        └─ virginica (1/1)
            └─ PetalLength < 4.85
                └─ SepalWidth < 3.1
                    └─ virginica (2/2)
                        └─ versicolor (1/1)
                └─ virginica (43/43)
```

Confusion Matrix and Overfitting

It is common practice to summarize the performance of a model using a *confusion matrix*, containing the true positives and negatives found by our model on the test data, as well as the false positives and negatives.

In the case of binary classification, a confusion matrix is shaped as follows.

Predicted / Ground truth	Positive	Negative
Positive	TP	FN
Negative	FP	TN

Among the many, three important measures can be obtained by the matrix above: accuracy, precision, and recall. In the binary classification scenario, they are defined as follows.

$$\text{Accuracy} = \frac{TP + TN}{TP + FP + TN + FN}$$

$$\text{Precision} = \frac{TP}{TP + FP}$$

$$\text{Recall} = \frac{TP}{TP + FN}$$

In the multi-class scenario, as in our case, we can compute precision and recall individually for each class. For obtaining a unique scalar, we can average all the results.

```
In [16]: cm = confusion_matrix(y_predict, y)
```

```
Out[16]:
```

Predicted	Ground Truth		
	setosa	versicol...	virginica
setosa	50	0	0
versicol...	0	50	0
virginica	0	0	50

```
In [17]: # wow! our model is so good!
accuracy(cm)
```

```
Out[17]: 1.0
```

How awful! The model we just trained is bad, for sure.

Can you tell why?

Answer (decode from [base64encode](#)):

```
VGhlIGNvZGUgaXMgbm90IGdlbmVyYWxpemluZyEKVGhIHNwbGl0cyBpbIB0aGUgdHJ·
```

Model Evaluation

Imagine projecting the data points on a bidimensional plane: can you provide a graphical sketch of what is happening during the inference process of the tree trained above?

Let us to obtain a more reliable model.

```
In [18]: (X_train, X_test), (y_train, y_test) = partition((X, y), 0.7, rng=121, sh
```

```
In [19]: mach = machine(model, X_train, y_train)
fit!(mach)
y_predict_probabilities = MLJ.predict(mach, X_test)
y_predict = mode.(y_predict_probabilities)
cm = confusion_matrix(y_predict, y_test)
```

[Info: Training machine(DecisionTreeClassifier(max_depth = 5, ...), ...).

```
Out[19]:
```

Predicted	Ground Truth		
	setosa	versicol...	virginica
setosa	14	0	0
versicol...	0	13	1
virginica	0	2	15

We can iterate the process above on multiple *folds*, to assess the overall quality of a machine learning training strategy. This technique is commonly called *cross-*

validation.

In the following, the iris dataset will be shuffled and divided into training and test in different ways, and each time a decision tree will be learned and tested over a different portion of the data.

```
In [20]: acc = evaluate!(  
    mach,  
    resampling=StratifiedCV(; nfolds = 5, shuffle=true),      # cross valid  
    measures=[accuracy]  
)
```

Evaluating over 5 folds: 100%[=====] Time: 0:00:02

Out[20]: PerformanceEvaluation object with these fields:

```
model, measure, operation,  
measurement, per_fold, per_observation,  
fitted_params_per_fold, report_per_fold,  
train_test_rows, resampling, repeats
```

Extract:

measure	operation	measurement
Accuracy()	predict_mode	0.943

per_fold	1.96*SE
[0.952, 0.857, 1.0, 0.952, 0.952]	0.0511

Training with Hyperparameters Tuning

The arguments of `DecisionTreeClassifier(...)` are said to be **hyperparameters**, as they are the meta-parameters exploited for creating a specific algorithm (i.e., the if-else cascade we call decision tree).

Which combination of hyperparameters should we provide?

In this rather lightweight example, we can systematically try many combinations and keep the one which expresses the highest performance.

This technique goes under the name of *grid search*.

```
In [21]: max_depth_range = range(Int, :max_depth, lower=2, upper=10)  
min_samples_leaf_range = range(Int, :min_samples_leaf, lower=1, upper=5)  
min_samples_split_range = range(Int, :min_samples_split, lower=2, upper=1)
```

```
In [22]: tuned_tree = TunedModel(  
    model = MLJDecisionTreeInterface.DecisionTreeClassifier(),  
    resampling = StratifiedCV(nfolds = 5, shuffle = true),  
    range = [max_depth_range, min_samples_leaf_range, min_samples_split_r  
    measure = accuracy,  
    tuning = Grid()  
)
```

```
Out[22]: ProbabilisticTunedModel(  
    model = DecisionTreeClassifier(  
        max_depth = -1,  
        min_samples_leaf = 1,  
        min_samples_split = 2,  
        min_purity_increase = 0.0,  
        n_subfeatures = 0,  
        post_prune = false,  
        merge_purity_threshold = 1.0,  
        display_depth = 5,  
        feature_importance = :impurity,  
        rng = TaskLocalRNG()),  
    tuning = Grid(  
        goal = nothing,  
        resolution = 10,  
        shuffle = true,  
        rng = TaskLocalRNG()),  
    resampling = StratifiedCV(  
        nfolds = 5,  
        shuffle = true,  
        rng = TaskLocalRNG()),  
    measure = Accuracy(),  
    weights = nothing,  
    class_weights = nothing,  
    operation = nothing,  
    range = MLJBase.NumericRange{Int64, MLJBase.Bounded, Symbol}[NumericRa  
nge(2 ≤ max_depth ≤ 10; origin=6.0, unit=4.0), NumericRange(1 ≤ min_samp  
les_leaf ≤ 5; origin=3.0, unit=2.0), NumericRange(2 ≤ min_samples_split  
≤ 10; origin=6.0, unit=4.0)],  
    selection_heuristic = MLJTuning.NaiveSelection(nothing),  
    train_best = true,  
    repeats = 1,  
    n = nothing,  
    acceleration = CPU1{Nothing}(nothing),  
    acceleration_resampling = CPU1{Nothing}(nothing),  
    check_measure = true,  
    cache = true,  
    compact_history = true,  
    logger = nothing)
```

```
In [23]: # find the best model, exploring different hyperparameterizations leverag  
# cross validation  
mach = machine(tuned_tree, X, y)  
fit!(mach)  
y_predict_probabilities = MLJ.predict(mach, X_test)  
y_predict = mode.(y_predict_probabilities)  
cm = confusion_matrix(y_predict, y_test)  
  
[ Info: Training machine(ProbabilisticTunedModel(model = DecisionTreeClass  
ifier(max_depth = -1, ...), ...), ...).  
[ Info: Attempting to evaluate 405 models.  
Evaluating over 405 metamodels: 100%[=====] Time: 0:0  
0:00
```

Out[23]:

Predicted	Ground Truth		
	setosa	versicol...	virginica
setosa	14	0	0
versicol...	0	13	0
virginica	0	2	16

Training Forests

In [24]:

```
try
    RandomForestClassifier = @load RandomForestClassifier pkg=DecisionTree
catch
    println("The RandomForestClassifier symbol has already been imported.
end

import MLJDecisionTreeInterface ✓
[ Info: For silent loading, specify `verbosity=0`.
```

Out[24]: MLJDecisionTreeInterface.RandomForestClassifier

In [25]:

```
forest = MLJDecisionTreeInterface.RandomForestClassifier(
    max_depth=3, min_samples_leaf=1, min_samples_split=2, n_trees=10)

forestmach = machine(forest, X_train, y_train)

MLJ.fit!(forestmach, verbosity=0)

y_predict_probabilities = MLJ.predict(forestmach, X_test)
y_predict = mode.(y_predict_probabilities)

cm = confusion_matrix(y_predict, y_test)
```

Out[25]:

Predicted	Ground Truth		
	setosa	versicol...	virginica
setosa	14	0	0
versicol...	0	12	1
virginica	0	3	15

Learning with Sole.jl

Tabular Datasets and Logisets

Symbolic AI treats tabular datasets, such as the iris flower, as sets of propositional interpretations, onto which formulas of propositional logic are interpreted.

Look at the (classical) tabular dataset \mathcal{I} below. We denote instances with I , and

variables^[1], as V_i .

	V_1	V_2	V_3
I_1	1.2	[1, 2, 3]	A
I_2	1.3	[9, 7, 6]	B
I_3	0.8	[2, 8, 2]	C
I_4	1.1	[1, 3, 7]	B
I_5	1.2	[4, 3, 3]	B

We can change the point of view on the table above from a statistical to a logical one, called a *logiset*.

This requires the definition of a propositional alphabet \mathcal{P} .

Consider $\mathcal{P} = \{p, q, r\}$, with:

$$p \text{ coloneqq } V_1 \geq 1$$

$$q \text{ coloneqq } \text{sum}(V_2) < 13$$

$$r \text{ coloneqq } V_3 = \text{B}$$

We denote the truth constant with \top (top), and the false constant with \perp (bot).

The resulting (propositional) logiset $\mathcal{I}_{\mathcal{P}}$ is:

	p	q	r
I_1	\top	\top	\perp
I_2	\top	\perp	\top
I_3	\perp	\top	\perp
I_4	\top	\top	\top
I_5	\top	\top	\top

[1] We use the term "variable" to denote, in general, a column of the tabular dataset: this corresponds to a raw attribute or a *feature* (a processed attribute).

```
In [26]: using MLJBase
using SoleData
```

```
In [27]: X_logiset = PropositionalLogiset(data);
X_logiset.tabulardataset == data
```

```
Out[27]: true
```

```
In [28]: phi = parseformula(
    "SepalLength > 5.8 ∧ SepalWidth < 3.0 ∨ Species == \"setosa\"";
    atom_parser = a->Atom(
        parsecondition(
            SoleData.ScalarCondition, a;
            featuretype = SoleData.VariableValue
        )
    ),
)
```

```

)

```

Γ **Warning:** Please, specify a type for the feature values (featvaltype = ...). Real will be used, but note that this may raise type errors. (expr = "SepalLength > 5.8")
 ↳ @ SoleData ~/.julia/packages/SoleData/N5aE8/src/scalar/conditions.jl:330
 Γ **Warning:** Please, specify a type for the feature values (featvaltype = ...). Real will be used, but note that this may raise type errors. (expr = "SepalWidth < 3.0")
 ↳ @ SoleData ~/.julia/packages/SoleData/N5aE8/src/scalar/conditions.jl:330
 Γ **Warning:** Please, specify a type for the feature values (featvaltype = ...). Real will be used, but note that this may raise type errors. (expr = "Species == \"setosa\"")
 ↳ @ SoleData ~/.julia/packages/SoleData/N5aE8/src/scalar/conditions.jl:330

Out[28]: SyntaxBranch: (SepalLength > 5.8 ∧ SepalWidth < 3.0) ∨ Species == setosa

```
In [29]: # check(phi, SoleLogics.LogicalInstance(X_logiset, 1))
check(phi, X_logiset, 1)
```

Out[29]: true

From DecisionTree.jl to SoleModels.jl

If we manage to make an existing model compliant with the interface of `SoleModels` package, then we can play with it from a logical standpoint.

```
In [30]: using SoleModels
```

```
In [31]: mach = machine(model, X_train, y_train)
fit!(mach)

# \:seedling:
 = fitted_params(mach).tree
```

[**Info:** Training machine(DecisionTreeClassifier(max_depth = 5, ...), ...).

```
Out[31]: PetalLength < 2.6
    └─ setosa (36/36)
        └─ PetalWidth < 1.65
            └─ PetalLength < 4.95
                └─ versicolor (34/34)
                    └─ SepalLength < 6.15
                        └─ SepalWidth < 2.45
                            └─ virginica (1/1)
                                └─ versicolor (1/1)
                                    └─ virginica (2/2)
                            └─ virginica (31/31)
```

```
In [32]: # we encode the model in such a way that it can be investigated via SoleM
# \:evergreen_tree:
 = solemodel(
printmodel(
```

```

█ ([PetalLength] < 2.5999999999999996)
└✓ setosa
  └✗ ([PetalWidth] < 1.65)
    └✓ ([PetalLength] < 4.95)
      └✓ versicolor
        └✗ ([SepalLength] < 6.15)
          └✓ ([SepalWidth] < 2.45)
            └✓ virginica
              └✗ versicolor
                └✗ virginica
      └✗ virginica
  └✗ virginica

```

In [33]: `# these are all the logical rules encoded by the tree
listrules(iris)`

Out[33]: 6-element Vector{ClassificationRule{String}}:

- █ ([PetalLength] < 2.5999999999999996) ↳ setosa
- █ (([PetalLength] ≥ 2.5999999999999996)) ∧ (([PetalWidth] < 1.65)) ∧ (([PetalLength] < 4.95)) ↳ versicolor
- █ (([PetalLength] ≥ 2.5999999999999996)) ∧ (([PetalWidth] < 1.65)) ∧ (([PetalLength] ≥ 4.95)) ∧ (([SepalLength] < 6.15)) ∧ (([SepalWidth] < 2.45)) ↳ virginica
- █ (([PetalLength] ≥ 2.5999999999999996)) ∧ (([PetalWidth] < 1.65)) ∧ (([PetalLength] ≥ 4.95)) ∧ (([SepalLength] < 6.15)) ∧ (([SepalWidth] ≥ 2.45)) ↳ versicolor
- █ (([PetalLength] ≥ 2.5999999999999996)) ∧ (([PetalWidth] < 1.65)) ∧ (([PetalLength] ≥ 4.95)) ∧ (([SepalLength] ≥ 6.15)) ↳ virginica
- █ (([PetalLength] ≥ 2.5999999999999996)) ∧ (([PetalWidth] ≥ 1.65)) ↳ virginica

In [34]: `metricstable(iris)`

Antecedent	Consequent	ninstances	ncovered	coverage	confidence
lift	natoms				
[PetalLength] < 2.5999999999999996	setosa	105	36		
0.342857 1.0 2.91667 1					
	([PetalLength] ≥ 2.5999999999999996) ∧ ([PetalWidth] < 1.65) ∧ ([PetalLength] < 4.95) versicolor				
	105 34 0.32381 1.0 3.0 3				
	([PetalLength] ≥ 2.5999999999999996) ∧ ([PetalWidth] < 1.65) ∧ ([PetalLength] ≥ 4.95) ∧ ([SepalLength] < 6.15) ∧ ([SepalWidth] < 2.45) virginica				
	105 1 0.00952381 1.0 3.08824 5				
	([PetalLength] ≥ 2.5999999999999996) ∧ ([PetalWidth] < 1.65) ∧ ([PetalLength] ≥ 4.95) ∧ ([SepalLength] < 6.15) ∧ ([SepalWidth] ≥ 2.45) versicolor				
	105 1 0.00952381 1.0 3.0 5				
	([PetalLength] ≥ 2.5999999999999996) ∧ ([PetalWidth] < 1.65) ∧ ([PetalLength] ≥ 4.95) ∧ ([SepalLength] ≥ 6.15) virginica				
	105 2 0.0190476 1.0 3.08824 4				
	([PetalLength] ≥ 2.5999999999999996) ∧ ([PetalWidth] ≥ 1.65) virginica				
	105 31 0.295238 1.0 3.08824 2				

```
In [35]: # show all the testing instances to the tree, and compare the metrics
# with the testing samples
apply!(tree, X_test, y_test);
```

```
In [36]: # we can visualize how our model behaved at testing time
metricstable(
    tree;
    normalize = true,
    metrics_kwargs = (;),
    additional_metrics = (;),
    height = r->SoleLogics.height(antecedent(r))
)
)
```

Antecedent	Consequent	ninstances	ncovered	coverage	confidence
lift	natoms	height			
$h] < 2.5999999999999996$	setosa	45	14	0.311111	[PetalLength]
1.0 3.21429 1 0					$([PetalLength] \in [2.5999999999999996, 4.95))$
$\wedge ([PetalWidth] < 1.65)$	versicolor	45	13	0.288889	
1.0 3.0 2 1					$([PetalLength] \geq 4.95) \wedge ([PetalWidth] < 1.65) \wedge ([SepalLength] < 6.15)$
$\wedge ([SepalWidth] < 2.45)$	virginica	45	0	0.0	
NaN NaN 4 3					$([PetalLength] \geq 4.95) \wedge ([PetalWidth] < 1.65) \wedge ([SepalLength] < 6.15)$
$\wedge ([SepalWidth] \geq 2.45)$	versicolor	45	1	0.0222222	
0.0 0.0 4 3					$([PetalLength] \geq 4.95) \wedge ([PetalWidth] < 1.65) \wedge ([SepalLength] \geq 6.15)$
NaN NaN 3 2	virginica	45	0	0.0	
$\wedge ([PetalWidth] \geq 1.65)$	virginica	45	17	0.377778	$([PetalLength] \geq 2.5999999999999996)$
0.882353 2.48162 2 1					

```
In [37]: # join some rules for the same class into a single, sufficient and necessary
# condition for the same class
metricstable(joinrules(, min_ncovered = 1, normalize = true))
```

Antecedent	Consequent	ninstances	ncovered	coverage	confidence
lift	natoms	height			
$[PetalLength] < 2.5999999999999996$	setosa	45	14	0.311111	
0.311111 1.0 3.21429 1					$(([PetalLength] \in [2.5999999999999996, 4.95)) \wedge ([PetalWidth] < 1.65)) \vee$
$(([PetalLength] \geq 4.95) \wedge ([PetalWidth] < 1.65) \wedge ([SepalLength] < 6.15) \wedge$	versicolor	90	14	0.155556	$([SepalWidth] \geq 2.45))$
0.928571 2.78571 6					
$[PetalLength] \geq 2.5999999999999996 \wedge [PetalWidth] \geq 1.65$	virginica	45	17	0.377778	
45 17 0.377778 0.882353 2.48162 2					

Here, we are just scratching the surface of Sole framework, limiting ourselves to pretty printings.

In the next lessons, we will enhance the machine learning pipeline we introduced today, with spatial reasoning considerations.

Below, there is a little spoiler about a fancy machine learning model, which is general enough for dealing with more-than-propositional logics.

In [38]:

```
using ModalDecisionTrees

mdt_model = ModalDecisionTree()
mach = machine(mdt_model, X_test, y_test)
fit!(mach)
y_pred = predict_mode(mach)
cm = confusion_matrix(y_predict, y_test)
```

[Info: Precomputing logiset...

[Info: Training machine(ModalDecisionTree(max_depth = nothing, ...), ...).

Out[38]:

		Ground Truth		
		setosa	versicol...	virginica
Predicted	setosa	14	0	0
setosa	14	0	0	0
versicol...	0	12	1	1
virginica	0	3	15	15

Exercise

Try to write your own pipeline, considering the [seeds dataset](#).

In [39]:

```
using CSV
```

In [40]:

```
SEEDS_PATH = joinpath("../", "datasets", "seeds.csv")
data = DataFrame(
    CSV.File(SEEDS_PATH; header=false)
)
```

Out[40]: 210×8 DataFrame

185 rows omitted

Row	Column1	Column2	Column3	Column4	Column5	Column6	Column7	Co
	Float64	Int						
1	15.26	14.84	0.871	5.763	3.312	2.221	5.22	
2	14.88	14.57	0.8811	5.554	3.333	1.018	4.956	
3	14.29	14.09	0.905	5.291	3.337	2.699	4.825	
4	13.84	13.94	0.8955	5.324	3.379	2.259	4.805	
5	16.14	14.99	0.9034	5.658	3.562	1.355	5.175	
6	14.38	14.21	0.8951	5.386	3.312	2.462	4.956	
7	14.69	14.49	0.8799	5.563	3.259	3.586	5.219	
8	14.11	14.1	0.8911	5.42	3.302	2.7	5.0	
9	16.63	15.46	0.8747	6.053	3.465	2.04	5.877	
10	16.44	15.25	0.888	5.884	3.505	1.969	5.533	
11	15.26	14.85	0.8696	5.714	3.242	4.543	5.314	
12	14.03	14.16	0.8796	5.438	3.201	1.717	5.001	
13	13.89	14.02	0.888	5.439	3.199	3.986	4.738	
:	:	:	:	:	:	:	:	:
199	12.62	13.67	0.8481	5.41	2.911	3.306	5.231	
200	12.76	13.38	0.8964	5.073	3.155	2.828	4.83	
201	12.38	13.44	0.8609	5.219	2.989	5.472	5.045	
202	12.67	13.32	0.8977	4.984	3.135	2.3	4.745	
203	11.18	12.72	0.868	5.009	2.81	4.051	4.828	
204	12.7	13.41	0.8874	5.183	3.091	8.456	5.0	
205	12.37	13.47	0.8567	5.204	2.96	3.919	5.001	
206	12.19	13.2	0.8783	5.137	2.981	3.631	4.87	
207	11.23	12.88	0.8511	5.14	2.795	4.325	5.003	
208	13.2	13.66	0.8883	5.236	3.232	8.315	5.056	
209	11.84	13.21	0.8521	5.175	2.836	3.598	5.044	
210	12.3	13.34	0.8684	5.243	2.974	5.637	5.063	

In [41]: # write your pipeline here